

CASPO: Learned Prefix Value Modeling for Stepwise Credit Assignment in Reasoning RL

Anonymous

Abstract

We propose CASPO (Credit Assignment Step Policy Optimization), a method for reinforcement learning of mathematical reasoning that obtains step-level credit signals from a learned prefix-value model rather than from on-policy Monte Carlo prefix rollouts. CASPO retains the VinePPO step-segmentation and step-TD advantage structure but replaces VinePPO’s $K=9$ rollouts at each non-terminal prefix with a single forward pass through an IPVRM-style cumulative-log-ratio prefix value model, optionally fine-tuned online during policy training under an ADB+DLW correction. The result is a strict drop-in for VinePPO that uses the same PPO clipped objective, the same KL penalty, the same response segmentation, and the same global PPO minibatch, but eliminates the per-step Monte Carlo expansion that dominates VinePPO’s wall-clock cost. We compare four methods (GRPO, PPO+Critic, VinePPO, CASPO) under one trainer, one verifier, one rollout engine, and one set of hyperparameters, on Qwen2.5-Math-1.5B with the `dssr_sub` subset of DeepScaleR (1209 prompts from One-Shot-RLVR), with DeepSeekMath-7B on MATH-lighteval as a 7B paper-faithful reference. PPO+Critic is the canonical Schulman-2017 PPO with a learned scalar value head on a stripped-LM-head backbone copy and GAE; we use it rather than bare PPO (uniform sequence-broadcast advantages with no value baseline) because GRPO already covers the value-free case via its per-prompt group-relative baseline. PPO+Critic provides a per-token learned baseline that contrasts with CASPO’s per-step learned baseline and with VinePPO’s on-policy Monte-Carlo step values. Throughout we keep the algorithmic contract identical except for the source of the per-step value $V(s_t)$, isolating the credit-signal change from confounders. Empirical results are reported in Section 6 after the corresponding training and evaluation runs complete.

1 Introduction

Reinforcement learning has emerged as the dominant approach for improving the mathematical reasoning of large language models, but the dominant signal is sparse: a verifier returns 0/1 only at the end of a long chain of thought, leaving the optimizer to assign credit across hundreds of intermediate tokens. PPO [5] addresses this by broadcasting a sequence-level advantage uniformly across response tokens; GRPO [6] refines the variance term using a group-relative baseline. Neither method distinguishes a correct intermediate inference from a lucky guess that nonetheless reached the right final answer.

VinePPO [2] addresses this gap by introducing *step-level* value estimates: the response is segmented into reasoning steps, and the value at each step boundary s_t is estimated by sampling K on-policy continuations from s_t and averaging their terminal verifier rewards. This produces well-calibrated step values that bear out empirically—VinePPO reports substantial improvements over PPO and GRPO on MATH and GSM8K—but at a cost: each non-terminal step boundary requires K additional generations, multiplying wall-clock time by roughly the average step count per response.

This cost is most punishing precisely on the long-chain regimes where step-level credit is most needed. On the One-Shot-RLVR `dsr_sub` DeepScaleR subset with Qwen2.5-Math-1.5B, the LaTeX-aware step splitter produces ~ 40 step boundaries per response (vs. ~ 10 – 20 on upstream MATH), and the $K=9$ Monte Carlo continuations push per-step training to ~ 33 minutes on $4\times H100$ 80GB. A 600-outer-step VinePPO run at this configuration thus requires roughly 14 days of wall-clock and is practically infeasible to run as a baseline at our experimental scale; the cost-effectiveness gap to VinePPO is largest exactly where step-level credit assignment matters most.

CASPO replaces those Monte Carlo prefix rollouts with a single forward pass through a *learned* prefix value model $V_\phi(s_t)$. Following IPVRM [1], V_ϕ is parameterized as a cumulative log-ratio between a trainable prefix policy π_ϕ and a frozen reference π_{ref} , and trained from trajectory-level correctness labels using a margin BCE objective. CASPO keeps every other VinePPO component fixed: the same LaTeX-aware step splitter, the same step-TD advantage construction, the same clipped PPO objective, and the same KL penalty. Step-level credit is therefore extracted from V_ϕ in $\mathcal{O}(1)$ forward passes per response rather than $\mathcal{O}(K \cdot \text{steps})$ generations, and online updates of V_ϕ during policy training mitigate distribution shift.

Contributions. We contribute: (i) CASPO: a method that drops VinePPO’s prefix Monte Carlo step in favor of an IPVRM-style learned prefix value model, including the step-TD construction and an online update path with ADB and DLW corrections; (ii) a family of three principled step-TD advantage formulations (*value*, *prob*, *logprob*) derived from the Bradley–Terry interpretation of V_ϕ under the IPVRM BCE loss, each operating on a different scale of the prefix value (raw cumulative log-ratio, implied correctness probability, or its log) and motivated by a different bias–variance regime; (iii) a frozen reward-model variant (CASPO-frozen-RM) that disables online value updates and isolates the cost/benefit of online value drift correction; (iv) an iterated-refresh configuration that periodically retrains V_ϕ from scratch on Monte-Carlo-labeled rollouts of the current policy, converting the credit model from a single drifting model into a sequence of in-distribution snapshots, at amortized wall-clock cost well below one VinePPO outer-step budget; (v) a unified comparison stack that implements GRPO, PPO+Critic, VinePPO, and CASPO under one trainer, one verifier, and one rollout engine, controlling everything except the algorithmic credit-signal choice; and (vi) implementation infrastructure for full-finetune 7B-scale RL with FSDP-sharded trainer and rank-local vLLM generation, including a streaming CUDA-IPC weight-sync path that is collective with `summon_full_params`, fp32 master weights with on-demand AdamW CPU-offload during the sync window to fit the unsharded fp32 materialization, and per-method ckpt sharding across multiple NVMe drives so each of the four 1B methods runs to completion in parallel without disk contention.

Implementation provenance. The codebase is an adaptation and extension of the public VinePPO setup [2, 4]. The DeepSeekMath-7B configuration’s prompt template, rollout shape, PPO hyperparameters, evaluation protocol, and LaTeX-aware step segmentation are matched to or derived from the upstream paper and reference repository so that any wall-clock or accuracy gap between methods reflects the algorithmic change rather than infrastructure. The Qwen2.5-Math-1.5B configuration on `dsr_sub` adapts the same training contract to the model’s native prompt template and to a 4-GPU FSDP topology, and is the primary experimental site for the comparison and for the iterated-refresh contribution.

2 Related Work

Policy gradient with sparse rewards. PPO [5] optimizes a clipped policy-gradient surrogate to bound the per-step policy update. Applied to language models with verifier-only terminal rewards, the standard practice is to broadcast the sequence-level advantage uniformly across response tokens and add a KL-to-reference penalty to prevent the policy from drifting far from the SFT initialization.

Group-relative baselines. GRPO [6] samples G responses per prompt and uses the within-group reward statistics as the advantage baseline. This avoids training a separate value network entirely, at the cost of treating every intermediate token in a correct response as equally important.

Step-level credit assignment via Monte Carlo prefix values. VinePPO [2] introduces stepwise value estimates by sampling K on-policy continuations from each non-terminal prefix and averaging the verifier outcomes of those continuations. Step advantages are then formed as temporal-difference differences between consecutive prefix values. The reasoning-credit gain is well demonstrated on MATH, but the method’s rollout multiplier scales as K times the average number of steps per response.

Learned process supervision. A line of work trains a separate process reward model on human-annotated or self-supervised step-correctness labels. CASPO’s prefix value model is in this family but constructed differently: rather than scoring step text directly, it accumulates token-level $\log[\pi_\phi/\pi_{\text{ref}}]$ contributions, following IPVRM [1]. This aligns prefix scoring with the same density that generated the response and avoids requiring a separate textual annotation.

Implicit prefix value learning (IPVRM). IPVRM [1] formalizes prefix values as cumulative log-ratios of a trainable prefix policy versus a frozen reference policy, trained from trajectory-level outcome labels using a margin BCE. IPVRM’s reported setting uses LoRA adapters on a 7B base; we adopt the same parameterization but perform full-model updates, with corresponding adjustments to the online learning rate.

Per-token learned-critic baselines. The classical PPO-with-critic recipe [5] uses a learned scalar value head $V_\psi(s, t)$ trained jointly with the policy on the value loss $(V_\psi - \hat{R})^2$ and folded into the advantage via generalized advantage estimation (GAE). We include this recipe as a baseline (PPO+Critic) so that CASPO’s learned *step-level* prefix value can be compared not only against the value-free baselines (PPO, GRPO) and the on-policy Monte Carlo step value (VinePPO), but also against an in-line per-token learned baseline trained from the same trajectories under the same PPO objective.

Generation backends and weight sync. We use vLLM [3] for high-throughput rollout and prefix caching. The implementation uses CUDA-IPC weight transfer between trainer ranks and rank-local vLLM engines; for the FSDP-sharded 7B trainer, the IPC sync is implemented as a collective `summon_full_params` block over the FSDP-wrapped policy followed by streaming per-tensor IPC handle exchange with each rank’s vLLM `EngineCore`, in chunks of 8 parameters with per-parameter bfloat16 casting to bound transient host memory during the sync window.

3 Background

3.1 Reasoning RL setup

A trajectory consists of a problem prompt x and a model response $y = (y_1, \dots, y_T)$ generated autoregressively. A verifier $R(x, y) \in \{0, 1\}$ returns 1 if and only if the boxed final answer in y matches the ground truth. The training objective is to maximize $\mathbb{E}_{x \sim \mathcal{D}, y \sim \pi_\theta(\cdot|x)} R(x, y)$ subject to a KL constraint to a frozen reference policy π_{ref} , which we instantiate as the SFT initialization.

3.2 Step segmentation

We split each response into reasoning steps using the LaTeX-aware splitter of VinePPO [2, 4], which identifies natural mathematical step boundaries (display equations, aligned blocks, sentence terminators in prose). Let $0 = t_0 < t_1 < \dots < t_S = T$ be the resulting boundary token indices and let $s_t \triangleq y_{<t}$ denote the prefix state at boundary t .

3.3 Step-TD advantage

Given a per-step value function $V(s_t)$, both VinePPO and CASPO construct the step-TD advantage

$$A(s_t) = r_t + \gamma V(s_{t+1}) - V(s_t), \quad r_t = \begin{cases} R(x, y) & t = S, \\ 0 & t < S, \end{cases} \quad (1)$$

with $\gamma = 1$. CASPO obtains V from a learned model V_ϕ ; VinePPO obtains V by averaging K Monte Carlo continuations. The downstream PPO objective and KL penalty are identical for both methods; only the source of V differs.

3.4 PPO objective

We use the clipped PPO objective

$$\mathcal{L}_{\text{clip}}(\theta) = -\mathbb{E}_t[\min(\rho_t(\theta) A_t, \text{clip}(\rho_t(\theta), 1 - \epsilon, 1 + \epsilon) A_t)], \quad (2)$$

with $\rho_t(\theta) = \pi_\theta(y_t | x, y_{<t}) / \pi_{\theta_{\text{old}}}(y_t | x, y_{<t})$ and $\epsilon = 0.2$. A KL-to-reference term is added with the k_3 estimator:

$$\text{KL}_{k_3}(\pi_\theta \| \pi_{\text{ref}}) \approx \mathbb{E}_t[\exp(\delta_t) - \delta_t - 1], \quad \delta_t = \log \pi_{\text{ref}}(y_t) - \log \pi_\theta(y_t). \quad (3)$$

The advantages A_t in (2) are (a) terminal-broadcast group-relative advantages for GRPO; (b) GAE per-token advantages from a learned scalar value head for PPO+Critic; or (c) the step-TD advantages of (1) broadcast over the tokens in each step span for VinePPO/CASPO.

4 Method: CASPO

Figure 1 summarizes the comparison axis: rollout, step splitter, verifier, advantage form, PPO objective, and KL penalty are held identical across all four methods. Only the source of the per-step value $V(s_t)$ differs.




Figure 1: Four compared methods, anchored on the source of the per-step value $V(s_t)$. Rollout, segmenter, verifier, advantage form, PPO clipped objective, and KL penalty are identical across all rows; the only difference is which black box produces $V(s_t)$ used in $A_t = r_t + \gamma V(s_{t+1}) - V(s_t)$. GRPO uses no value ($V \equiv 0$) with a per-prompt group-relative terminal broadcast. PPO+Critic (the canonical Schulman-2017 PPO) learns a scalar per-token value head trained on $(V_\psi - \hat{R})^2$ with GAE. VinePPO estimates $V(s_t)$ by averaging $K=9$ on-policy continuations from each prefix. CASPO replaces those continuations with an IPVRM-style cumulative-log-ratio prefix value model V_ϕ .

4.1 Prefix value model

Following IPVRM [1], we define

$$V_\phi(s_t) = \beta \sum_{i < t} \log \frac{\pi_\phi(y_i | x, y_{<i})}{\pi_{\text{ref}}(y_i | x, y_{<i})}, \quad (4)$$

where π_ϕ is a trainable prefix policy initialized from the same SFT checkpoint family as the base policy and π_{ref} is the same frozen reference used in the KL term in (3). The scale β is a fixed scalar; we use $\beta = 10$ throughout.

4.2 Phase-1: offline value training

We train V_ϕ on a fixed corpus of (prompt, response, outcome) triples collected by sampling G responses per prompt from the SFT policy at generation temperature 1.0, scoring with the verifier, and discarding prompts whose G outcomes are all 0 or all 1 (the mixed-outcome filter, following IPVRM). For a response of length T we use the length-normalized prefix value $\bar{V}_\phi(s_t) := V_\phi(s_t)/t$ and minimize a per-prefix BCE-with-margin against the eventual binary outcome $r_o \in \{0, 1\}$ (IPVRM Eq. 9):

$$\mathcal{L}_{\text{IPVRM}}(\phi) = -\mathbb{E}_{(x, y, r_o)} \left[\frac{1}{T} \sum_{t=1}^T \left(r_o \log \sigma(\bar{V}_\phi(s_t) - m) + (1 - r_o) \log(1 - \sigma(\bar{V}_\phi(s_t) + m)) \right) \right], \quad (5)$$

with margin $m = 5$. We train for three epochs with full-model updates and a small learning rate (5×10^{-7}).

4.3 Phase-2: step-TD policy optimization

At each PPO outer step, the trainer samples G responses per prompt for P prompts, segments each response, and constructs step values $\{V_\phi(s_t)\}$ in a single forward pass through V_ϕ using (4). Step advantages are then assembled from (1) and broadcast over the tokens within each step span before applying the PPO objective in (2) with KL penalty in (3). The standard CASPO run optionally co-trains V_ϕ during policy optimization (next section).

4.4 Online value updates with ADB+DLW

Distribution drift between the SFT-init prefix policy and the actively optimized policy can degrade V_ϕ . We therefore allow online updates of V_ϕ from the same rollout batch used for PPO. The on-policy update uses the same margin-BCE loss (5) but with two corrections from IPVRM:

- *ADB* (advantage debias): the BCE boundary is shifted per-prompt by $\log[\mu/(1-\mu)]$, where μ is the empirical positive fraction over the prompt’s G rollouts, so the value model learns *deviations* from the prompt’s SFT difficulty rather than its absolute difficulty.
- *DLW* (difficulty-level weighting): the BCE term is weighted by $\mu(1-\mu)$, downweighting prompts where outcomes are nearly all-correct or all-incorrect.

The online value learning rate is set to 10^{-6} , an order of magnitude below the offline rate to control drift.

4.5 Three step-TD advantage formulations from V_ϕ

A central modeling choice in CASPO is *which* statistic of V_ϕ enters the step-TD difference. We derive and propose three principled formulations from how the IPVRM loss (5) parameterizes correctness. Recall that the BCE margin loss treats V_ϕ as a Bradley–Terry score: the implied per-prefix correctness probability is

$$\tilde{p}(s_t) \triangleq \sigma(V_\phi(s_t)) \approx \mathbb{P}[R(x, y) = 1 \mid s_t], \quad (6)$$

where σ is the logistic function. Equation (6) is the natural way to convert an unbounded cumulative-log-ratio score into a calibrated probability, and it is the quantity the BCE objective directly optimizes.

Equation (1) is then well-defined on three distinct scales of V_ϕ , each with a different statistical interpretation. We propose all three as named formulations of CASPO:

value (raw cumulative log-ratio):

$$A^V(s_t) = r_t + \gamma V_\phi(s_{t+1}) - V_\phi(s_t). \quad (7)$$

prob (Bradley–Terry probability):

$$A^P(s_t) = r_t + \gamma \tilde{p}(s_{t+1}) - \tilde{p}(s_t) = r_t + \gamma \sigma(V_\phi(s_{t+1})) - \sigma(V_\phi(s_t)). \quad (8)$$

logprob (log Bradley–Terry probability):

$$A^{LP}(s_t) = r_t + \gamma \log \tilde{p}(s_{t+1}) - \log \tilde{p}(s_t). \quad (9)$$

In all three variants, the terminal verifier reward $r_T = R(x, y)$ enters identically; the only difference is the function applied to V_ϕ before the TD difference.

Intuition. The three transforms have different scales and different bias–variance profiles:

- *value*: The TD operates on the unbounded log-ratio scale. A single token whose log-ratio swings by, say, -10 produces an unbounded ΔV_ϕ , which carries a strong gradient signal but also amplifies miscalibrated regions of the prefix value model (e.g. tokens far from the ref-policy distribution). High signal-to-noise where V_ϕ is well-calibrated, but heavier tails when it is not. Standardization + clipping (Section 4, advantage clip ± 3) bounds the magnitude empirically, but does not reshape the underlying scale.
- *prob*: The TD operates on the probability scale, so $|A^P(s_t)| \leq 1$ before standardization. By (6), $\tilde{p}(s_t) \approx \mathbb{E}[R \mid s_t]$, so A^P is the empirical advantage that an idealized binary-reward critic would compute: it is the change in the expected terminal reward attributed to step t . This is the variant most directly compatible with the binary verifier we use, but it suffers *near saturation*: when $V_\phi(s_t)$ is large positive or negative, $\sigma'(V_\phi)$ collapses toward zero, so two prefixes with very different log-ratios get essentially the same gradient. In practice this means *prob* undersignals on prefixes the value model is highly confident about.
- *logprob*: The TD operates on the log-probability scale, mirroring the form of token log-likelihoods used in the PPO ratio $\rho_t(\theta)$. For a sequence model trained by maximum likelihood, log-probability differences are the canonical update signal; using $\log \tilde{p}$ keeps the value-derived advantage on the same information-theoretic scale as the underlying token logits. Numerically, $\log \sigma(V) = -\log(1 + e^{-V})$, which behaves linearly for $V \ll 0$ (matching the value scale on the negative tail) and saturates softly to zero for $V \gg 0$ (compressing the high-confidence positive tail). This is a smooth interpolation between *value* (linear in the negative regime) and *prob* (bounded above).

When each formulation is preferable. The three formulations are not interchangeable: each is the principled choice under a different calibration regime of V_ϕ . *prob* ((8)) is the natural statistical match when V_ϕ is well-calibrated as a Bradley–Terry score, and admits a clean expected-terminal-reward interpretation $A^P(s_t) = r_t + \gamma \mathbb{E}[R \mid s_{t+1}] - \mathbb{E}[R \mid s_t]$ via (6). *value* ((7)) is preferred when V_ϕ has wide log-ratio dynamic range that the BCE objective does not directly bound, in which case operating on $\sigma(V_\phi)$ would discard signal at the saturation tails; the price is heavier-tailed advantages, controlled in our setup by ± 3 standardized clipping. *logprob* ((9)) provides a smooth interpolation: it preserves linear sensitivity in the negative- V_ϕ regime (where most gradient signal lives early in training, when the policy is still mostly off-distribution) and softly compresses the positive- V_ϕ regime (where over-confident prefixes would otherwise dominate the batch). It also matches the information-theoretic scale of the underlying token logits used in the PPO ratio $\rho_t(\theta)$, which is a desirable invariance property under reparameterization of π_ϕ . We treat these three as a contributed family of step-TD formulations and report all three on MATH-500 pass@16 in Section 6.

4.6 Step advantage normalization and clipping

Let $\mathcal{A} = \{A(s_t)\}$ over all valid step boundaries in the outer PPO step. We standardize over the entire batch ($\tilde{A}_t = (A_t - \mu_{\mathcal{A}})/(\sigma_{\mathcal{A}} + \varepsilon)$) and clip the standardized advantages to $[-3, 3]$ to bound the contribution of single-step outliers. The clip range is chosen empirically to balance suppression of pathological prefixes against retaining genuine step-level signal.




Figure 2: The three step-TD transforms of the prefix value V_ϕ . *value* (dashed) is the identity: linear sensitivity, unbounded tails. *prob* (solid) is bounded in $[0, 1]$ but saturates symmetrically, so prefixes the model is highly confident about contribute near-zero gradient. *logprob* (dotted) is the smooth interpolation: linear in the negative- V_ϕ regime where most early-training signal lives, soft saturation toward zero in the positive- V_ϕ regime that compresses overconfident prefixes. The TD difference $A^*(s_t) = r_t + \gamma f(V_\phi(s_{t+1})) - f(V_\phi(s_t))$ inherits the regime of f .

5 Experimental Setup

5.1 Models and data

The primary experimental site is **Qwen2.5-Math-1.5B** (Qwen/Qwen2.5-Math-1.5B), a 1.5B math-specialist SFT initialization, trained on the **One-Shot-RLVR** `dsr_sub` 1209-prompt subset of DeepScaleR. Evaluation uses MATH-500, GSM8K, and OlympiadBench under greedy decoding ($k=1$, $T=0$). Qwen2.5-Math is a math-specialist SFT model that responds best to its native template `{query}\nLet's think step by step and output the final answer within \boxed{}`.; substituting the `[MATH_TASK]` template degrades Qwen accuracy by >30 percentage points and is a critical reproduction footgun.

For a 7B paper-faithful reference site, we additionally evaluate on **DeepSeekMath-7B** (realtreetune/deepseek-7b-math), the SFT-on-MATH initialization from the upstream VinePPO pipeline, trained on DigitalLearningGmbH/MATH-1.1B (train split). 7B held-out evaluation uses HuggingFaceH4/MATH-500 (test split, 500 problems) plus the full MATH test set, College Math, and OlympiadBench (math-only configuration), under the upstream prompt template `[MATH_TASK] Problem:\n{query}\n\nSolution:` and the LaTeX-aware step splitter from the upstream codebase.

5.2 Methods

We compare four methods (Figure 1):

GRPO Per-prompt group-relative terminal-reward advantages with $G = 8$. Stands in for the value-free baseline; bare PPO (uniform sequence-broadcast advantages, no group baseline, no value head) is strictly weaker on this evaluation regime and is omitted.

PPO+Critic Schulman-2017 PPO with a learned scalar value head V_ψ on a stripped-LM-head backbone copy, trained jointly with the policy under a clipped value loss $\max[(V_\psi - \hat{R})^2, (\text{clip}(V_\psi, V_\psi^{\text{old}} \pm \epsilon_v) - \hat{R})^2]$ with $\epsilon_v = 0.2$ and $\lambda_{\text{GAE}} = 1.0$, value-loss coefficient 1.0, critic LR 10^{-6} .

Table 1: Training hyperparameters at the two experimental sites. The Qwen2.5-Math-1.5B configuration is the primary site for the comparison; DeepSeekMath-7B retains the upstream VinePPO configuration as a paper-faithful 7B reference. The 2048-token response budget for Qwen2.5-Math-1.5B is justified empirically (see text below).

Field	Qwen2.5-Math-1.5B	DeepSeekMath-7B
Group size G	8	8
Prompts per outer step	128	64
Responses per outer step	1024	512
Global PPO minibatch	128	64
PPO epochs per rollout	2 (GRPO reported at both $\mu=1$ and $\mu=2^a$)	2
Policy LR	1×10^{-6}	1×10^{-6}
Online value LR (CASPO)	1×10^{-6}	1×10^{-6}
Critic LR (PPO+Critic)	1×10^{-6}	1×10^{-6}
Value β , m	10, 5	10, 5
PPO clip ϵ	0.2	0.2
Value clip ϵ_v (PPO+Critic)	0.2	0.2
GAE λ (PPO+Critic)	1.0	1.0
KL coefficient	10^{-3} (k_3 , 10^{-2} for PPO+Critic)	10^{-2} (k_3)
Advantage clip	3.0	3.0
Advantage standardization	batch	batch
Rollout sampling	temp 0.6, top- p 1.0	temp 0.6, top- p 0.9
Max prompt / response	1024 / 2048	1512 / 1024
Max sequence length	3072	2499
Warmup steps	0	480
Outer steps	600	1000
Checkpoint cadence	every 50	every 200

^a PPO+Critic, VinePPO, and CASPO are fixed at $\mu=2$ epochs per rollout, matching the upstream VinePPO configuration. GRPO is reported at both $\mu=1$ (canonical setting in DeepSeekMath [6], TRL, and verl) and $\mu=2$ (matches the optimizer-step budget per rollout of the other three methods, following the VinePPO-paper GRPO baseline). The $\mu=1$ comparison is GRPO-faithful; the $\mu=2$ comparison isolates the credit-signal change from optimizer-budget effects.

VinePPO $K = 9$ Monte Carlo continuations from each non-terminal step boundary; step-TD advantages.

CASPO IPVRM prefix value model with online updates; step-TD advantages over V_ϕ . We report all three step-TD formulations from Section 4.5 (CASPO-*value*, CASPO-*prob*, CASPO-*logprob*) as separate runs, and a fourth variant CASPO-frozen-RM that uses the *value* formulation but disables online updates of V_ϕ during policy training.

5.3 Hyperparameters

All four methods share the same PPO infrastructure. Table 1 summarizes the key knobs.

The Qwen2.5-Math-1.5B configuration uses 1024 responses per outer step (128 prompts $\times G=8$) on a 4-GPU FSDP topology. The KL coefficient of 10^{-3} for CASPO/GRPO/VinePPO (versus 10^{-2} for PPO+Critic) is empirically required: the math-specialist Qwen SFT initialization is over-anchored at 10^{-2} and policy updates do not move the model without a weaker reference anchor,

while PPO+Critic is critic-stabilized and tolerates the stronger anchor. The DeepSeekMath-7B configuration retains the upstream VinePPO hyperparameters at 7B [2, 4]: 512 responses per outer step ($64 \times G=8$), 2 PPO epochs per rollout, and a 64-response global PPO minibatch, isolating any empirical differences between the two sites to the model and dataset rather than the optimization recipe.

The response budget of 2048 tokens for Qwen2.5-Math-1.5B is set by empirical measurement on `dsr_sub`: an 800-chain uncapped sample ($T=1.0$, $\text{max}=3000$) places the 98th percentile of *correct* chains at 1613 tokens. At 2048, only 2% of correct chains are truncated, while the overall truncation rate is 8% (i.e., the truncated tail is overwhelmingly off-track / rambling chains rather than length-bound correct ones). This makes the unfinished-response penalty signal load the right pattern: “long without resolution” is what gets penalized, not “hard problem requires more tokens.” Going further to 3072 catches fewer than one additional correct chain per 150 sampled and risks the Qwen `max_position_embeddings=4096` ceiling at prompt + response.

5.4 Evaluation protocol

Final evaluation reports `pass@1` and `pass@k` ($k = 16$) on MATH-500, full MATH test, College Math, and OlympiadBench. Sampling at evaluation uses temperature 0.35 and top- p 0.9, matching the upstream protocol. Generation is performed with vLLM at `kv_cache_dtype="fp8"` and `gpu_memory_utilization=0.92` for maximum eval throughput on a dedicated GPU. Sample evaluation at intermediate checkpoints uses `EVAL_LIMIT=100`, `EVAL_K=8` for low-overhead training-curve visibility; the full benchmark suite is run at the final checkpoint.

5.5 Infrastructure

Trainer + colocated rollout. Both scales share a single PyTorch trainer plus rank-local vLLM rollout. The 1B configuration uses a single H100 per method (Figure 3, left); the 7B configuration uses `FSDP=4` (`full_shard`, `use_orig_params`) for GRPO/PPO+Critic/CASPO and `FSDP=8` for VinePPO, with vLLM colocated rank-local in both layouts (Figure 3, right).

Streaming CUDA-IPC weight sync. At 1B with a single-rank trainer, the IPC handle exchange is direct. At 7B under FSDP, the trainer enters `FSDP.summon_full_params(writeback=False)` collectively, then streams parameters to each rank’s vLLM `EngineCore.collective_rpc("update_weights", ...)` in chunks of 8 tensors with per-parameter bf16 cast inside the IPC push. Streaming bounds peak host-side staging memory and avoids materializing the entire unsharded fp32 policy on host before transfer.

fp32 master weights with on-demand AdamW offload. Both methods co-train policy (and CASPO’s value model, or PPO+Critic’s critic) in bf16 compute with fp32 master weights. Under FSDP, the master copy lives in the AdamW optimizer state and is materialized into bf16 shards through `MixedPrecision(param_dtype=bf16, reduce_dtype=fp32)`. At sync time the `summon_full_params` block briefly materializes a ~ 28 GiB unsharded fp32 buffer per rank for the 7B policy. To fit alongside the already-resident ~ 14 GiB of AdamW state and the ~ 24 GiB vLLM KV cache reservation per H100, the trainer offloads the AdamW state to pinned host memory before entering the `summon` block and restores it after IPC push, costing ~ 400 – 500 ms per sync at PCIe Gen5 but freeing the GPU memory required for the unsharded fp32 materialization.

1B fp32 master via autocast. At 1B (single-GPU, no FSDP) the trainer wraps each forward in `torch.autocast(dtype=bfloat16, cache_enabled=False)` so the fp32 parameters cast to bf16 for compute on every op. The non-cached cast is required for AdamW optimizer-state coverage to include all parameters; with caching enabled, the autocast layer holds bf16 tensor wrappers that are not written through to the underlying fp32 parameters during the optimizer step.

Multi-drive checkpoint sharding for parallel 1B training. Each method’s 1B run produces $\sim 15\text{--}29$ GB per checkpoint (model in bf16, plus the AdamW master+momentum bundles for policy, plus either CASPO’s value-model AdamW or PPO+Critic’s critic AdamW). At `save_every=100` over 1000 outer steps that is roughly 150–290 GB per method-run. To run all four 1B method-runs (GRPO, PPO+Critic, CASPO, CASPO-*prob*) concurrently on four GPUs of the same node at the `save_every=100` cadence, each method’s outputs are routed to a dedicated 350 GB NVMe volume, so the four runs neither compete for disk bandwidth nor cross-fill each other’s drives.

Reference-policy sharing. The frozen reference policy is shared between the KL term and CASPO’s value-model ref to avoid a duplicate forward. Rollout, old-logprob rescore, ref-logprob precompute, and policy/value forward all use FlashAttention-3 on Hopper.


Paper-faithful alignment with upstream VinePPO. A targeted audit against the upstream VinePPO reference repository [4] produced a set of small but consequential alignment fixes that we include in the comparison stack so that our VinePPO baseline reproduces the upstream numerics: prepending the BOS token at both rollout and training, per-microbatch advantage whitening, adding the prompt-style stop string `"\n\n\nProblem:"` for vLLM, clamping per-token KL into $[0, 10]$ before reduction, penalizing unfinished sequences via the max-sequence-length policy in (2), an explicit PPO ratio safety skip at $10\times$, and uniform grad-accumulation per outer step. With these in place, the VinePPO baseline in our stack matches the upstream configuration field-for-field at the respective scale.

Stabilization deviations from upstream (PPO+Critic/CASPO only). Two deliberate departures from upstream-faithful numerics were required to keep the online-value methods (PPO+Critic, CASPO, CASPO-p) stable at 1B. Empirically, with upstream’s `kl_coef =  $10^{-4}$` and no advantage clip, both PPO+Critic and CASPO drift to $\text{KL} \rightarrow \infty$ within ~ 20 outer steps in our reimplementation; GRPO remains stable. The fix was a stronger ref-anchor and an outlier safety net:


- `kl_coef =  $10^{-2}$` ($100\times$ upstream’s 10^{-4} , midway between InstructGPT’s 0.02 and DeepSeek-Math GRPO’s 0.04). This is roughly the standard RLHF range; upstream VinePPO’s unusually small value is paired with $K=9$ Monte Carlo step values that absorb V-noise we do not have.
- Standardized advantage clip $\pm 3\sigma$ retained as an outlier safety net (the upstream-faithful no-clip choice survives at VinePPO but causes CASPO to occasionally take large policy steps that destabilize online V_ϕ).

GRPO and VinePPO are unaffected by these knobs at runtime since they have no learned baseline to drift; we hold the same values across methods for stack consistency.

(a) 1B per method (single H100, four methods $\rightarrow$ four GPUs)



(b) 7B FSDP=4 (or 8 for VinePPO) + rank-local colocated vLLM



streaming IPC push, chunks of 8 with bf16 cast | AdamW state CPU-offloaded across the `summon` window

Figure 3: Trainer/rollout topology. (a) Each 1B run colocates trainer and vLLM on a single H100 (four methods $\rightarrow$ four GPUs of the same node); fp32 master is maintained via `autocast(cache_enabled=False)` so AdamW covers all parameters, and the IPC handle exchange between the single-rank trainer and its vLLM engine is direct. (b) Each 7B run uses FSDP=4 (FSDP=8 for VinePPO) with one rank-local vLLM engine per rank. Trainer-to-vLLM weight sync streams parameters in chunks of 8 with per-tensor bf16 cast inside an `FSDP.summon_full_params(writeback=False)` block; AdamW state is offloaded to pinned host memory across the `summon` window so the ~ 28 GiB unsharded fp32 materialization fits alongside the vLLM KV cache reservation.

V_ϕ retrain after verifier and tokenization changes. V_ϕ 's offline training data was collected before two upstream-aligned patches (Minerva-style answer-extraction cascade in the verifier; explicit BOS prepending in the rollout engine). At RL time, this leaves V_ϕ scoring *prefixes that begin with a different first token* (`<s>` vs. `[]`) against *outcomes from a verifier that accepts a broader set of answer formats*. The offline-trained V_ϕ reaches step-1 `v_acc`= 0.32 against runtime, well below its offline-train accuracy of ~ 0.96 . Re-running the data-collection pipeline with the live verifier and rollout engine, then retraining V_ϕ for three epochs at the same paper-faithful 5×10^{-7} offline learning rate, lifts step-1 `v_acc` to 0.74 at runtime — a $2.3\times$ recovery from the stale-label/shifted-prompt mismatch. We use this retrained V_ϕ for all reported CASPO runs.

6 Results

This section will report wall-clock and accuracy after the full training and evaluation runs complete. The intended structure is:

- Per-method MATH-500 pass@1 and pass@16 at **final**, with seed-level error bars over a small number of seeds (target: 3 seeds per method per scale). Methods include all four compared rows: GRPO, PPO+Critic, VinePPO, and the three CASPO step-TD formulations plus CASPO-frozen-RM.

- Per-method full-suite (MATH, College Math, OlympiadBench) pass@16 at **final**.
- Training curves on MATH-500 sample evaluation at intermediate checkpoints (1B: every 100 outer steps; 7B: every 200 outer steps) plus **final**, summarizing sample efficiency.
- Per-formulation comparison of the three step-TD advantage variants (CASPO-*value* vs. CASPO-*prob* vs. CASPO-*logprob*) on MATH-500 pass@16, mapping the empirical winner to the calibration regime predicted in Section 4.5 and visualized in Figure 2.
- CASPO-frozen-RM vs. CASPO (default, online value updates on) to quantify the contribution of the online update path.
- PPO+Critic vs. CASPO: contrast a per-token learned scalar baseline against a per-step learned prefix value, both under the same PPO objective and the same global PPO minibatch.
- Wall-clock summary: per-method step time, per-method 1000-step ETA, and the implied GPU-hour cost of the four-method comparison, with the breakdown by phase (rollout, old-logprob rescore, value/critic forward, ref-logprob precompute, policy update, IPC sync).

We report all numbers under the same training contract: identical SFT initialization, identical prompt set, identical verifier, identical group/response budget, identical PPO clipped objective, identical KL coefficient, identical advantage standardization, identical max sequence length, and identical evaluation sampling parameters. The only intended difference between rows is the source of the per-step value $V(s_t)$ in (1): zero with a per-prompt group baseline (GRPO), a learned per-token scalar value $V_\psi(s, t)$ (PPO+Critic), Monte Carlo continuations (VinePPO), or a learned prefix value $V_\phi(s_t)$ (CASPO).

7 Analysis

This section will examine, given the empirical results, what determines when V_ϕ matches VinePPO’s prefix-rollout signal and when it does not. The intended structure is:

- *Step-credit calibration*: do step advantages from V_ϕ rank correct vs. incorrect step continuations consistently with on-policy MC?
- *Online value drift*: does the online update preserve the phase-1 calibration, or does it collapse to constant predictions? We monitor $\bar{V}^+$ and $\bar{V}^-$ (mean prefix value at correct vs. incorrect step continuations) over outer steps to detect drift.
- *Wall-clock per accuracy point*: the practical question CASPO is asked to answer is whether learned prefix values can match VinePPO’s accuracy at a fraction of the wall-clock cost. We report this as accuracy-versus-GPU-hours for each method.
- *Step-count effects on VinePPO*: VinePPO’s wall-clock scales as $K \cdot \text{steps/response}$, so the marginal cost of step-level credit grows with response length. We quantify this by binning the test set into long vs. short reasoning regimes and reporting the per-bin accuracy/wall-clock trade-off.

8 Limitations

Reward-model drift. CASPO’s online value updates are gated by ADB and DLW corrections, but no correction perfectly defends against pathological feedback loops (e.g. a value head that learns to predict the policy’s surface form rather than its correctness). The CASPO-frozen-RM variant in our contributed family isolates this risk by holding V_ϕ fixed during policy training, at the cost of distribution-shift bias.

A complementary direction is *periodic refresh from scratch*: at intervals of k outer steps the prefix value model is re-trained on a fresh Monte-Carlo-labeled corpus drawn from the current policy and a base- SFT initialization, then loaded as the new V_ϕ for the next k outer steps. This converts the credit model into a sequence of in-distribution snapshots rather than a single drifting model. The Qwen2.5-Math-1.5B configuration uses this iterated-refresh recipe (alternating phase length $k=150$ for the Δp transform, $k=200$ for $\Delta \log p$); the per-refresh wall-clock cost is approximately one-tenth of one VinePPO outer-step budget at $K=9$, leaving the method dominated by RL rollout cost rather than refresh overhead. We treat refresh-PRM as a distinct configuration of CASPO rather than a separate method, since it changes only how V_ϕ is sourced and not the step-TD advantage construction. A complete cross-method comparison of refresh-PRM CASPO against fixed- V_ϕ CASPO and VinePPO on Qwen2.5-Math-1.5B is in progress and is the primary subject of the companion experimental log accompanying this manuscript.

Cross-distribution probe-set evaluation of V_ϕ . Selecting the best refresh recipe by ranking candidate V_ϕ checkpoints on a fixed probe set is sensitive to the response-length cap used when the probe data was generated. A V_ϕ trained on collection budget L_{tr} sees prefixes spanning the range $[0, L_{\text{tr}}]$, but a probe set generated at cap $L_{\text{probe}} < L_{\text{tr}}$ contains prefixes only in $[0, L_{\text{probe}}]$. Differences in Spearman ρ between checkpoints trained at different L_{tr} on a single shared probe set therefore conflate intrinsic credit-assignment quality with the fraction of the checkpoint’s training distribution that the probe never exercises. We encountered this confound in early refresh-recipe sweeps and recommend that any future V_ϕ recipe selection generate the probe at the maximum training cap of any candidate, or otherwise restrict the comparison to recipes whose training caps match the probe cap.

Segmenter quality. The LaTeX-aware step splitter is a structural heuristic. Manual spot checks on Qwen2.5-Math-1.5B and DeepSeekMath-7B generations show coherent boundaries for prose, equations, factorization, and aligned derivations, but generations that continue past a final boxed answer have those trailing fragments treated as additional steps; a post-answer truncation pass is left as future work.

Apples-to-apples scope. Our comparison fixes the SFT initialization, prompt set, verifier, sampling parameters, KL constraint, and advantage standardization. It does not exhaustively explore method-specific tuning (e.g. different K for VinePPO, different β/m for CASPO, larger G for GRPO); we report the upstream-paper values where available and otherwise the values that minimize wall-clock without regressing pass@1 on a small validation set.

Scale. Results at 1.5B and 7B may not transfer to substantially larger scales. The 7B FSDP+vLLM-IPC stack is novel to this work; whether the same configuration extrapolates to 70B is an open systems question.

9 Conclusion

We described CASPO, a step-level credit assignment method for reasoning RL that replaces VinePPO’s on-policy Monte Carlo prefix rollouts with an IPVRM-style learned prefix value model and an optional online value update. CASPO fits cleanly into the same PPO clipped objective, the same KL constraint, and the same step-TD advantage construction used by VinePPO, isolating the credit-signal change from confounders. By unifying GRPO, PPO+Critic, VinePPO, and CASPO under one trainer, one verifier, and one rollout engine, we make it possible to compare these four methods under identical conditions on Qwen2.5-Math-1.5B with `dsr_sub` as the primary site and DeepSeekMath-7B with MATH-lighteval as a paper-faithful 7B reference. Empirical results in Section 6 will quantify whether the learned value model recovers VinePPO’s reasoning-credit signal at a fraction of its generation cost, or whether step-level Monte Carlo remains uniquely advantageous, and whether iterated refresh of V_ϕ from current-policy rollouts closes any residual gap.

References

- [1] Shiping Gao, Hongzhan Chen, Xiaojun Quan, Qifan Wang, and Lifu Huang. Unleashing implicit rewards: Prefix-value learning for distribution-level optimization. *arXiv preprint arXiv:2604.13197*, 2026. URL <https://arxiv.org/abs/2604.13197>.
- [2] Amirhossein Kazemnejad, Milad Aghajohari, Eva Portelance, Alessandro Sordoni, Siva Reddy, Aaron Courville, and Nicolas Le Roux. Vineppo: Unlocking rl potential for llm reasoning through refined credit assignment. *arXiv preprint arXiv:2410.01679*, 2024. URL <https://arxiv.org/abs/2410.01679>.
- [3] Woosuk Kwon, Zhuohan Li, Siyuan Zhuang, Ying Sheng, Lianmin Zheng, Cody Hao Yu, Joseph E. Gonzalez, Hao Zhang, and Ion Stoica. Efficient memory management for large language model serving with pagedattention. In *Proceedings of the ACM SIGOPS Symposium on Operating Systems Principles*, 2023. URL <https://arxiv.org/abs/2309.06180>.
- [4] McGill NLP. Vineppo codebase. <https://github.com/McGill-NLP/VinePPO>, 2024. Code for VinePPO: Unlocking RL Potential For LLM Reasoning Through Refined Credit Assignment.
- [5] John Schulman, Filip Wolski, Prafulla Dhariwal, Alec Radford, and Oleg Klimov. Proximal policy optimization algorithms. *arXiv preprint arXiv:1707.06347*, 2017. URL <https://arxiv.org/abs/1707.06347>.
- [6] Zhihong Shao, Peiyi Wang, Qihao Zhu, Runxin Xu, Junxiao Song, Xiao Bi, Haowei Zhang, Mingchuan Zhang, Y. K. Li, Y. Wu, Daya Guo, et al. Deepseekmath: Pushing the limits of mathematical reasoning in open language models. *arXiv preprint arXiv:2402.03300*, 2024. URL <https://arxiv.org/abs/2402.03300>.